

Streaming Algorithms

Meng-Tsung Tsai

05/01/2018

References

- "On Graph Problems in a Semi-Streaming Model," Feigenbaum et al. (2005)

Matching

Problem Definition

Input: a sequence of edges e_1, e_2, \dots, e_m of an n -node m -edge undirected graph G .

Output: a matching M of G .

Goal: make M as large as possible using $O(n \text{ poly log } n)$ space and few passes.

1/2-approximation, on a RAM

```
Matching(G){  
  M ← ∅;  
  foreach(e){  
    if(M ∪ {e} is a matching){  
      M ← M ∪ {e};  
    }  
  }  
  return M;  
}
```

$|M| \geq (1/2)\text{OPT}$. (Why?)

2/3-approximation, on a RAM

```
Matching(G){  
  M ← ∅;  
  foreach(e){  
    if(M ∪ {e} is a matching){  
      M ← M ∪ {e};  
    }  
  }  
  while(there exists a length-3 augmenting path P w.r.t. M){  
    M ← M ⊕ P;  
  }  
  return M;  
}
```

$|M| \geq (2/3)\text{OPT}$. (Why?)

Exercise 1

Devise an analogous algorithm that yields a $(1-\epsilon)$ -approximation for matching for any constant $\epsilon > 0$.

$(2/3-\epsilon)$ -approximation for bipartite graphs in the semi-streaming model

Augmentable length-3 paths

Let X be a maximum-size set of simultaneously augmentable length-3 paths with respect to M .

$$|M| + |X| \geq (2/3)\text{OPT. (Why?)}$$

Augmentable length-3 paths

Let X be a maximum-size set of simultaneously augmentable length-3 paths with respect to M .

$$|M| + |X| \geq (2/3)\text{OPT. (Why?)}$$

Let X' be a maximal-size set of simultaneously augmentable length-3 paths with respect to M .

$$|X'| \geq (1/3)|X|. \text{ (Why?)}$$

Semi-streaming Algorithm (First Attempt)

Matching(G) {

$M \leftarrow$ any maximal matching of G ;

 for($r = 1$; $r \leq 1/(2\epsilon)$; $++r$) {
 $X_r \leftarrow$ any $(1/3)$ -apx of X ;
 $M \leftarrow M \oplus X_r$;
 }

 return M ;

Matching(G) needs $O(1/(2\epsilon))$ passes if X can be $(1/3)$ -approximated in $O(1)$ passes, after which $|M| \geq (2/3 - \epsilon)\text{OPT}$. (Why?)

Semi-streaming Algorithm (Second Attempt)

Matching(G) {

$M \leftarrow$ any maximal matching of G ;

 for($r = 1$; $r \leq \left\lceil \frac{\log 6\epsilon}{\log 8/9} \right\rceil$; $++r$) {
 $X_r \leftarrow$ any $((1 - 2\delta)/3)$ -apx of X ; // $\delta = \epsilon/(2 - 3\epsilon)$
 $M \leftarrow M \oplus X_r$;
 }

 return M ;

Matching(G) needs $O(\log(1/\epsilon)/\epsilon)$ passes if X can be $((1 - 2\delta)/3)$ -approximated in $O(1/\delta)$ passes, after which $|M| \geq (2/3 - \epsilon)\text{OPT}$. (See the calculation in Ref. 1.)

$((1-2\delta)/3)$ -apx of X

Approximating- $X(M, \delta)$ { // order all edges in M from one partite to the other

```

V_C ← ∅; // used nodes
P_3 ← ∅; // length-3 augmenting paths
while(1){
  P_2 ← ∅; // length-2 augmenting paths
  foreach(edge (x, u) ∈ E/M){
    if(∃ ordered (u, v) ∈ M and x, u, v ∉ V_C)
      V_C ← V_C ∪ {x, u, v}, P_2 ← P_2 ∪ {(x, u, v)};
  }
  if (|P_2| ≤ δ|M|) return P_3; // |X| ≤ 2δ|M|
  foreach(edge (v, y) ∈ E/M){
    if(∃ ordered (x, u, v) ∈ P_2 and x, u, v, y ∉ V_C)
      V_C ← V_C ∪ {x, u, v, y}, P_3 ← P_3 ∪ {(x, u, v, y)};
  }
}
return P_3;

```

$O(2\delta)$ passes because each round except the last marks at least $\delta|M|$ nodes, and each round uses 2 passes.

$((1-2\delta)/3)$ -apx of X

Approximating- $X(M, \delta)$ { // order all edges in M from one partite to the other

```

V_C ← ∅; // used nodes
P_3 ← ∅; // length-3 augmenting paths
while(1){
  P_2 ← ∅; // length-2 augmenting paths
  foreach(edge (x, u) ∈ E/M){
    if(∃ ordered (u, v) ∈ M and x, u, v ∉ V_C)
      V_C ← V_C ∪ {x, u, v}, P_2 ← P_2 ∪ {(x, u, v)};
  }
  if (|P_2| ≤ δ|M|) return P_3; // |X| ≤ 2δ|M|
  foreach(edge (v, y) ∈ E/M){
    if(∃ ordered (x, u, v) ∈ P_2 and x, u, v, y ∉ V_C)
      V_C ← V_C ∪ {x, u, v, y}, P_3 ← P_3 ∪ {(x, u, v, y)};
  }
}
return P_3;

```

Every length-3 path in P_3 kills 3 paths in X .

$((1-2\delta)/3)$ -apx of X

Approximating- $X(M, \delta)$ { // order all edges in M from one partite to the other

```

V_C ← ∅; // used nodes
P_3 ← ∅; // length-3 augmenting paths
while(1){
  P_2 ← ∅; // length-2 augmenting paths
  foreach(edge (x, u) ∈ E/M){
    if(∃ ordered (u, v) ∈ M and x, u, v ∉ V_C)
      V_C ← V_C ∪ {x, u, v}, P_2 ← P_2 ∪ {(x, u, v)};
  }
  if (|P_2| ≤ δ|M|) return P_3; // |X| ≤ 2δ|M|
  foreach(edge (v, y) ∈ E/M){
    if(∃ ordered (x, u, v) ∈ P_2 and x, u, v, y ∉ V_C)
      V_C ← V_C ∪ {x, u, v, y}, P_3 ← P_3 ∪ {(x, u, v, y)};
  }
}
return P_3;

```

Every length-2 path in P_2 kills 2 paths in X .