

# Streaming Algorithms

Meng-Tsung Tsai


05/22/2018

## References

- "Communication Complexity," Kushilevitz
- "Lower Bounds for One-Way Communication," Roughgarden

## Randomized Communication Complexity

### The indexing problem



How many bits are needed to be sent from Alice to Bob so that Bob can figure out whether  $A[k] = 1$  or not w.p.  $> 7/8$  for every single input? ( $\Omega(n)$  bits)

## Distributional Complexity

Let  $P$  be a distribution over  $\{0, 1\}^n \times \{0, 1\}^n$ ; that is,  $P$  describes the probability that input =  $(x, y)$  for each  $(x, y) \in \{0, 1\}^n \times \{0, 1\}^n$ .  
Let  $f$  be the function to be computed.

If every deterministic one-way protocol  $D$  with

$$\Pr_{(x, y) \sim P}[D(x, y) \neq f(x, y)] \leq \epsilon$$

requires at least  $k$  bits, then every (public coin) randomized one-way protocol  $R$  with (two-sided) error  $\leq \epsilon$  requires at least  $k$  bits.

## Distributional Complexity

Proof.

Let  $R$  be a randomized protocol, i.e. a distribution over deterministic protocols  $D_1(R), D_2(R), \dots, D_s(R)$ .

If  $R$  requires  $< k$  bits to answer with error  $\leq \epsilon$ , then each deterministic protocol  $D_i(R)$  uses  $< k$  bits. By the assumption, each  $D_i(R)$  answers with error  $> \epsilon$ .

Then  $R$  answers with error  $> \epsilon$ .

## Deterministic Protocol that Allows Errors

Claim. If a deterministic protocol  $D$  for Index problem sends at most  $cn$  bits ( $c$  is a sufficiently small constant and  $n$  is sufficiently large) and the input is sampled uniformly, then the probability that  $D$  incurs an error is at least  $1/8$ .

Proof. For each message  $\mathbf{z}$  that Alice sends to Bob, depending only on  $\mathbf{z}$  and  $k$ , Bob has to answer  $A[k] = 0$  or  $1$ . Let  $b(\mathbf{z}) \in \{0, 1\}^n$  be Bob's answer when the transmitted message is  $\mathbf{z}$ .

Let  $S(\mathbf{z})$  be the set of Alice's input so that the transmitted message is  $\mathbf{z}$ .

- At most one element in  $S(\mathbf{z})$  has 0 disagreement with  $b(\mathbf{z})$ .
- At most  $n$  elements in  $S(\mathbf{z})$  have only 1 disagreement with  $b(\mathbf{z})$ .
- At most  $C(n, k)$  elements in  $S(\mathbf{z})$  have exactly  $k$  disagreement with  $b(\mathbf{z})$ .

## Deterministic Protocol that Allows Errors

Proof. There are  $2^{cn}$  distinct messages  $\mathbf{z}$ , and # elements in  $S(\mathbf{z})$  for all  $\mathbf{z}$  that have an error w.p.  $\leq n/4$  is at most

$$2^{cn} \left( \binom{n}{0} + \binom{n}{1} + \dots + \binom{n}{n/4} \right) \leq n 2^{cn} (4e)^{n/4} \leq 2^{(c+0.87)n + \log n}$$

Pick  $c < 0.03$  and  $n > 100$  so that  $((c+0.87)n + \log n) < n-1$ .

Hence, at least  $2^{n-1}$  (half) distinct inputs incur an error w.p.  $\geq 1/4$ .

For sufficiently large  $n$ , any  $D$  that uses  $0.03n$  bits has an error rate at least  $1/8$ .

$\Rightarrow$  For sufficiently large  $n$ , any  $R$  that has an error rate  $< 1/8$  for every single input requires  $\Omega(n)$  bits. QED

## Allowing Higher Error Rates

If  $R(\text{Index}) = o(n)$  when the randomized protocol  $R$  incurs an error w.p.  $< 1/2 - \epsilon$  for an arbitrary small constant  $\epsilon > 0$ , then we run a constant copies of  $R$  independently in parallel still requires  $o(n)$  bits.

On the other hand, the error rate drops to  $< 1/8$ .  $\rightarrow \leftarrow$

$$R(\text{Index}) = \Omega(n).$$

## Space Lower Bounds for Randomized Streaming Algorithms

## Topological Sort

Input: a directed acyclic graph  $D$

Output: a node ordering  $v_1, v_2, \dots, v_n$  so that for every arc  $(u, v)$  in  $D$ , node  $u$  appears earlier than  $v$  in the ordering.

Goal: show that any 1-pass randomized streaming algorithm that can output the ordering w.p.  $> 1/2 + \epsilon$  for any constant  $\epsilon > 0$  requires  $\Omega(n^2)$  bits.

Remark. The naive algorithm that stores the entire graph in an adjacency matrix turns out to be optimal.

## Reduction

Conduct a reduction from  $\text{Index}(C(n, 2))$  to T-sort.

Construct a graph with  $2n$  nodes,  $1x, 2x, \dots, nx, 1y, 2y, \dots, ny$ , initially without any edge.

Look at the problem instance of the Index problem. If the  $(i^*n+j)$ -th bit is 0, then add an edge from  $ix$  to  $jy$ , or otherwise add an edge from  $jx$  to  $iy$ .



## Reduction

To figure whether the  $(i*n+j)$ -th bit is 0 or 1, we add two edges  $(iy, ix)$  and  $(jy, jx)$ .

Claim. The resulting graph is still acyclic. (Why?)

If node  $iy$  appears before node  $jx$ , then the  $(i*n+j)$ -th bit is 0; otherwise  $iy$  appears after node  $jx$ , then the  $(i*n+j)$ -th bit is 1.



## Sorting

Input:  $n$  integers.

Output: the input integers in the sorted order.

Approach: Let Alice's input be an array of length  $2n$  and let half of them be 1-bits. In such cases,  $R(\text{Sorting})$  still has a lower bound  $\Omega(n)$ . Then we reduce Index to Sorting.

$$R(\text{Sorting}) = \Omega(n).$$

## Closest Pair

Input:  $n$  points on a plane.

Output: the pair of points whose distance is shortest.

$$R(\text{ClosestPair}) = \Omega(n).$$